

# **Mask Region-Based Convolutional Neural Networks for the Detection and Measurement of Copper Chalcogenide Nanoparticles in Transmission Electron Microscope Images**

## **Team 6**

Authors: Colin Dingley, Alex Davis

Course: AI 570 - Deep Learning

Summer September, 2024

### **▼ Project Background Information:**

The goal of this project is to develop a Mask R-CNN model for analyzing TEM images of nanoparticles. The original images were obtained from Dr. Kate Plass at Franklin and Marshall College. The aim is to automate the detection and measurement of Chalcogenide Nanoparticles.

This code was generated after consulting the work of bnsreenu and their github repository: [https://github.com/bnsreenu/python\\_for\\_microscopists/blob/master/330\\_Detectron2\\_Instance\\_3D\\_EM\\_Platelet.ipynb](https://github.com/bnsreenu/python_for_microscopists/blob/master/330_Detectron2_Instance_3D_EM_Platelet.ipynb)

### **▼ Installation and Setup**

### **▼ Install Libraries and Mount Google Drive**

All image and annotation files are stored on google drive. The following mounts the users google drive to CoLab to allow for access as though the files were local files.

```
import sys, os, distutils.core

# Install Detectron2
!git clone 'https://github.com/facebookresearch/detectron2'
dist = distutils.core.run_setup("./detectron2/setup.py")
!python -m pip install {' '.join([f"'{x}'" for x in dist.install_requires])}
sys.path.insert(0, os.path.abspath('./detectron2'))

# Install Tensorboard
%pip install tensorboard
%load_ext tensorboard

# Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')

# Setting Current Working Directory to Project Folder
prog_folder = "/content/drive/MyDrive/Colab_Notebooks/AI570_Project"
os.chdir(prog_folder)
print(os.getcwd())

↳ Cloning into 'detectron2'...
remote: Enumerating objects: 15743, done.
remote: Total 15743 (delta 0), reused 0 (delta 0), pack-reused 15743
Receiving objects: 100% (15743/15743), 6.32 MiB | 10.47 MiB/s, done.
Resolving deltas: 100% (11492/11492), done.
Ignoring dataclasses: markers 'python_version < "3.7"' don't match your environment
Requirement already satisfied: Pillow>=7.1 in /usr/local/lib/python3.10/dist-packages (9.4.0)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.7.1)
Requirement already satisfied: pycocotools>=2.0.2 in /usr/local/lib/python3.10/dist-packages (2.0.8)
Requirement already satisfied: termcolor>=1.1 in /usr/local/lib/python3.10/dist-packages (2.4.0)
Collecting yacs>=0.1.8
    Downloading yacs-0.1.8-py3-none-any.whl.metadata (639 bytes)
Requirement already satisfied: tabulate in /usr/local/lib/python3.10/dist-packages (0.9.0)
Requirement already satisfied:云pickle in /usr/local/lib/python3.10/dist-packages (2.2.1)
```

```
Requirement already satisfied: tqdm>=4.29.0 in /usr/local/lib/python3.10/dist-packages (4.66.4)
Requirement already satisfied: tensorboard in /usr/local/lib/python3.10/dist-packages (2.15.2)
Collecting fvcore<0.1.6,>=0.1.5
  Downloading fvcore-0.1.5.post20221221.tar.gz (50 kB)
    ━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 50.2/50.2 kB 3.2 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Collecting iopath<0.1.10,>=0.1.7
  Downloading iopath-0.1.9-py3-none-any.whl.metadata (370 bytes)
Collecting omegaconf<2.4,>=2.1
  Downloading omegaconf-2.3.0-py3-none-any.whl.metadata (3.9 kB)
Collecting hydra-core>=1.1
  Downloading hydra_core-1.3.2-py3-none-any.whl.metadata (5.5 kB)
Collecting black
  Downloading black-24.4.2-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (77 kB)
    ━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 77.1/77.1 kB 5.0 MB/s eta 0:00:00
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (24.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (4.53)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.4.5)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (1.26.4)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib) (2.)
Requirement already satisfied: PyYAML in /usr/local/lib/python3.10/dist-packages (from yacs>=0.1.8) (6.0.1)
Requirement already satisfied: absl-py>=0.4 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (1.4.0)
Requirement already satisfied: grpcio>=1.48.2 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (1.64.1)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (1.25.0)
Requirement already satisfied: google-auth-oauthlib<2,>=0.5 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (0.5.2)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (3.6)
Requirement already satisfied: protobuf!=4.24.0,>=3.19.6 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (3.21.2)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (2.29.0)
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (71.3.1)
Requirement already satisfied: six>1.9 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (1.16.0)
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (0.7.0)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard) (3.0.3)
Collecting portalocker (from iopath<0.1.10,>=0.1.7)
  Downloading portalocker-2.10.1-py3-none-any.whl.metadata (8.5 kB)
Collecting antlr4-python3-runtime==4.9.* (from omegaconf<2.4,>=2.1)
  Downloading antlr4-python3-runtime-4.9.3.tar.gz (117 kB)
    ━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 117.0/117.0 kB 11.5 MB/s eta 0:00:00
  Preparing metadata (setup.py) ... done
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from black) (8.1.7)
```

```
Collecting mypy-extensions>=0.4.3 (from black)
  Downloading mypy_extensions-1.0.0-py3-none-any.whl.metadata (1.1 kB)
```

## ▼ Importing libraries and functions

```
import torch, detectron2
import numpy as np
import os, json, cv2, random
from google.colab.patches import cv2_imshow
import matplotlib.pyplot as plt
import torchvision
import pandas as pd
from detectron2.utils.logger import setup_logger
from detectron2 import model_zoo
from detectron2.engine import DefaultPredictor, DefaultTrainer
from detectron2.config import get_cfg
from detectron2.utils.visualizer import Visualizer, ColorMode
from detectron2.data import MetadataCatalog, DatasetCatalog
from detectron2.data import build_detection_train_loader, build_detection_test_loader
from detectron2.data.datasets import register_coco_instances
from detectron2.evaluation import COCOEvaluator, inference_on_dataset
from detectron2.utils.events import TensorboardXWriter, EventStorage
from torch.utils.tensorboard import SummaryWriter

writer = SummaryWriter("./output/logs")
setup_logger()
```

```
<Logger detectron2 (DEBUG)>
```

## ▼ Detectron2 Evaluation

Detectron2 was developed by Facebook and was trained on a broad range of images. It should do well on detecting, segmenting and identifying regular images. To demonstrate this, we will run a pre-trained model on the following image of a street. The image was selected based on a google image search for non-copyrighted images of a street with people. We will use the pretrained model to evaluate its ability to analyse everyday images.

```
test_image = cv2.imread("./data/images/street_image.jpg")
scaled_image = cv2.resize(test_image, (0,0), fy=0.4, fx=0.4)
cv2_imshow(scaled_image)
```



```
def visualize_predictions(image_path, predictor, scale=0.7):
    """
    Visualize predictions on a random sample of images from the dataset.

    Args:
        image_path (str): Path to the image file.
        predictor: The Detectron2 predictor object.
        scale(float): Scale factor for the visualizer. Default is 0.7.
    """

    # Read and scale image
    image = cv2.imread(image_path)
    scaled_image = cv2.resize(image, (0,0), fy=scale, fx=scale)
```

```
# Make predictions on scaled image
predictions = predictor(scaled_image)

# Draw the predictions on the image.
labels = Visualizer(scaled_image[:, :, ::-1],
                     MetadataCatalog.get(cfg.DATASETS.TRAIN[0]))
                     #instance_mode=ColorMode.IMAGE_BW)
labeled_image = labels.draw_instance_predictions(predictions["instances"].to("cpu"))

# Combine original and labeled images side by side
combined_image = np.hstack((scaled_image, labeled_image.get_image()[:, :, ::-1]))

# Add titles
title_height = 50
title_image = np.zeros((title_height, combined_image.shape[1], 3), dtype=np.uint8)
cv2.putText(title_image, "Original",
            (scaled_image.shape[1] // 2 - 40, title_height // 2 + 10),
            cv2.FONT_HERSHEY_SIMPLEX, 1, (255, 255, 255), 2)
cv2.putText(title_image, "Labeled",
            (scaled_image.shape[1] + scaled_image.shape[1] // 2 - 40, title_height // 2 + 10),
            cv2.FONT_HERSHEY_SIMPLEX, 1, (255, 255, 255), 2)

combined_image_with_title = np.vstack((title_image, combined_image))

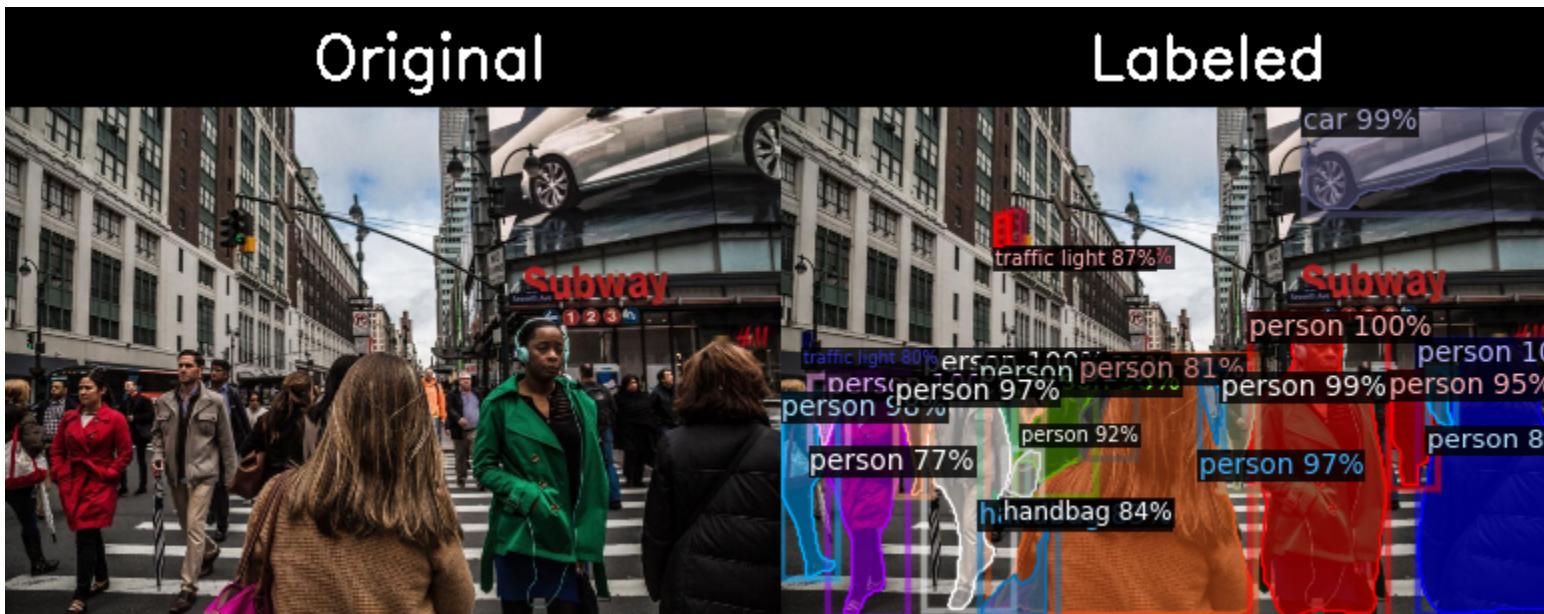
# Display the combined image with titles
cv2.imshow(combined_image_with_title)
```

We will be using facebook's detectron2's "model zoo" model. Details of this model can be found here: # [https://github.com/facebookresearch/detectron2/blob/main/MODEL\\_ZOO.md](https://github.com/facebookresearch/detectron2/blob/main/MODEL_ZOO.md)

```
cfg = get_cfg()
cfg.merge_from_file(model_zoo.get_config_file("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml"))
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.6 # set threshold for this model
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml")
predictor = DefaultPredictor(cfg)
```

```
outputs = predictor(scaled_image)
[07/30 14:43:21 d2.checkpoint.detection_checkpoint]: [DetectionCheckpointer] Loading from https://dl.fbaipublicfiles.com
model_final_f10217.pkl: 178MB [00:01, 90.3MB/s]
/usr/local/lib/python3.10/dist-packages/torch/functional.py:512: UserWarning: torch.meshgrid: in an upcoming release
    return _VF.meshgrid(tensors, **kwargs) # type: ignore[attr-defined]

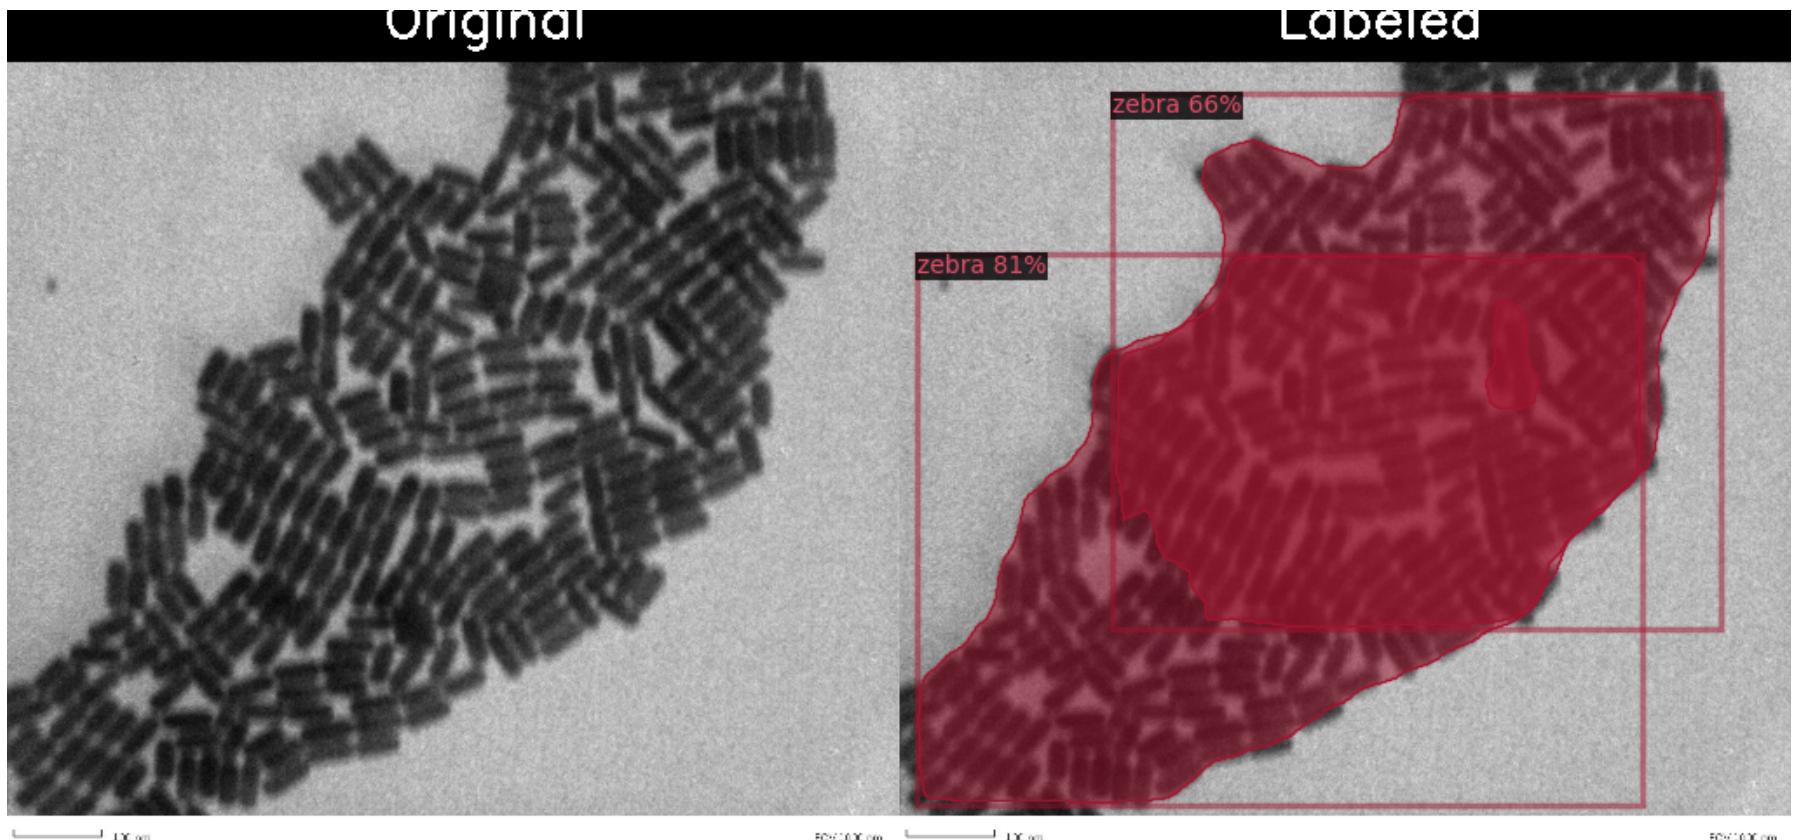
image_path = "./data/images/street_image.jpg"
visualize_predictions(image_path, predictor, scale=0.4)
```



## ▼ Default Model on TEM Images

The model did very well on identifying everyday objects, which the model was trained on. However, the model was not trained on TEM images. Let's see how well it performs on one of our images:

```
image_path = "./data/images/all/AE-002_000004.jpg"
visualize_predictions(image_path, predictor, scale=0.4)
```



Since zebras cannot fit into a TEM the labels are clearly incorrect. We will therefore need to refit the last layer of the model using our own labeled images in order to take advantage of transfer learning.

- ▼ Training on our labeled TEM images

```
def clear_dataset():
    """
```

```
....  
Checks for the existence of the dataset and removes it if it exists.  
This is needed for rerunning the notebook, or if the dataset is altered.  
"""
```

```
if "my_dataset_train" in MetadataCatalog.list():  
    MetadataCatalog.remove("my_dataset_train")  
    DatasetCatalog.remove("my_dataset_train")  
  
if "my_dataset_val" in MetadataCatalog.list():  
    MetadataCatalog.remove("my_dataset_val")  
    DatasetCatalog.remove("my_dataset_val")  
  
if "my_dataset_test" in MetadataCatalog.list():  
    MetadataCatalog.remove("my_dataset_test")  
    DatasetCatalog.remove("my_dataset_test")
```

## ▼ Loading and Registering COCO Datasets

```
clear_dataset()  
# Register datasets  
register_coco_instances("my_dataset_train", {}, "./data/annotations/train_annotations.json", "./data/images/all/")  
register_coco_instances("my_dataset_val", {}, "./data/annotations/val_annotations.json", "./data/images/all/")  
register_coco_instances("my_dataset_test", {}, "./data/annotations/test_annotations.json", "./data/images/all/")  
  
# Retrieve metadata and dictionaries for the training, testing and validation  
# datasets from the DatasetCatalog  
train_metadata = MetadataCatalog.get("my_dataset_train")  
train_dataset_dicts = DatasetCatalog.get("my_dataset_train")  
val_metadata = MetadataCatalog.get("my_dataset_val")  
val_dataset_dicts = DatasetCatalog.get("my_dataset_val")  
test_metadata = MetadataCatalog.get("my_dataset_test")  
test_dataset_dicts = DatasetCatalog.get("my_dataset_test")  
  
WARNING [07/30 14:44:03 d2.data.datasets.coco]:  
Category ids in annotations are not in [1, #categories]! We'll apply a mapping for you.
```

```
[07/30 14:44:03 d2.data.datasets.coco]: Loaded 14 images in COCO format from ./data/annotations/train_annotations.json
WARNING [07/30 14:44:04 d2.data.datasets.coco]:
Category ids in annotations are not in [1, #categories]! We'll apply a mapping for you.

[07/30 14:44:04 d2.data.datasets.coco]: Loaded 21 images in COCO format from ./data/annotations/val_annotations.json
WARNING [07/30 14:44:04 d2.data.datasets.coco]:
Category ids in annotations are not in [1, #categories]! We'll apply a mapping for you.

[07/30 14:44:04 d2.data.datasets.coco]: Loaded 11 images in COCO format from ./data/annotations/test_annotations.json
```

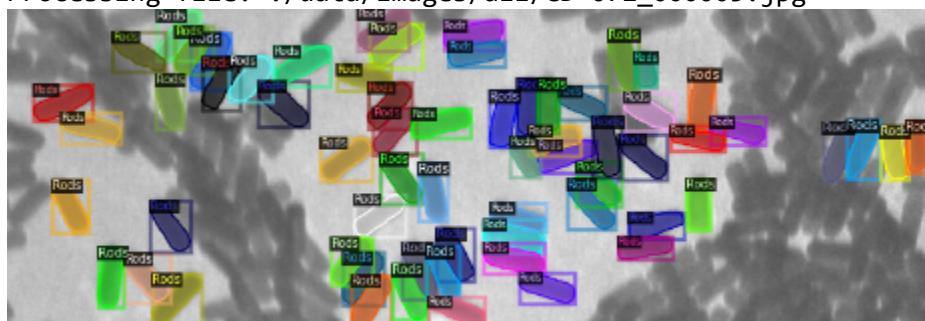
## ✓ Visualize Some Random Samples

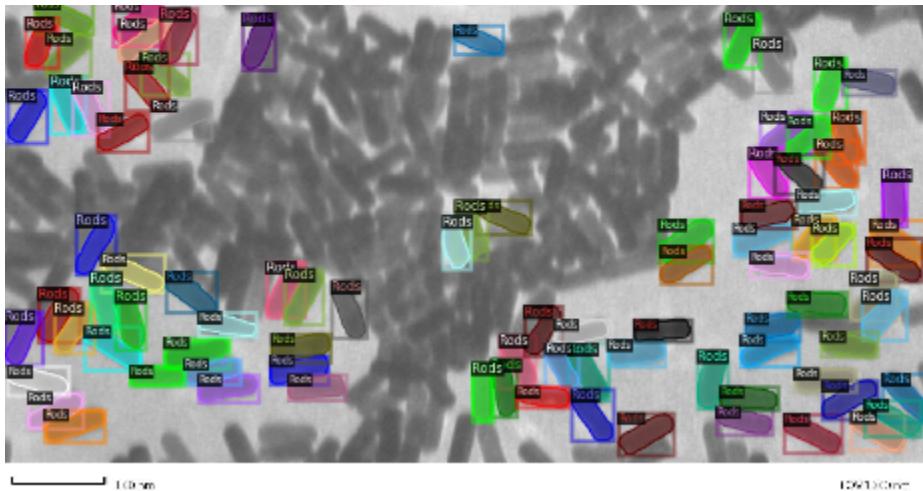
```
# Visualize some random samples
for d in random.sample(train_dataset_dicts, 2):
    print(f"Processing file: {d['file_name']}")
    image = cv2.imread(d["file_name"])
    if image is None:
        print(f"Error loading image {d['file_name']}")
        continue

    visualizer = Visualizer(image[:, :, ::-1], metadata=train_metadata, scale=0.6)
    vis = visualizer.draw_dataset_dict(d)

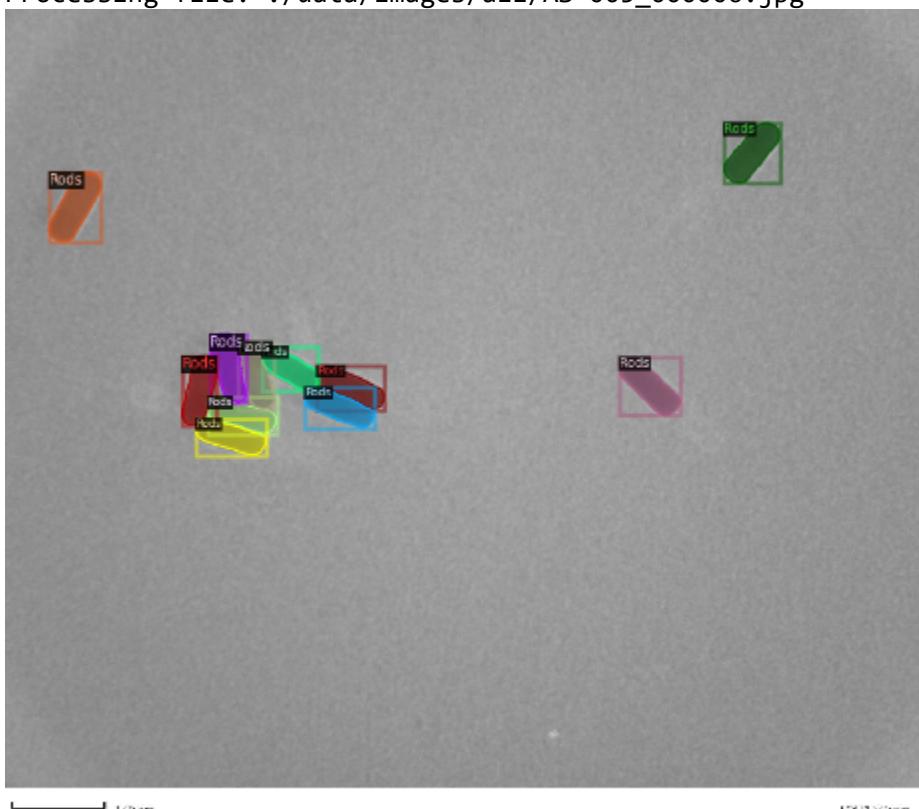
# Display the image with annotations
labeled_image = vis.get_image()[:, :, ::-1]
scaled_labeled_image = cv2.resize(labeled_image, (0,0), fy=0.6, fx=0.6)
cv2_imshow(scaled_labeled_image)
```

Processing file: ./data/images/all/CD-072\_000005.jpg





Processing file: ./data/images/all/AS-005\_000006.jpg



## ▼ Train

### ▼ Setting and configuring the model

```
class TrainerWithTensorboard(DefaultTrainer):
    @classmethod
    def build_writers(cls, cfg):
        return [TensorboardXWriter(cfg.OUTPUT_DIR)]

cfg = get_cfg()
cfg.OUTPUT_DIR = "./models/Detectron2_Models"
cfg.merge_from_file(model_zoo.get_config_file("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("my_dataset_train",)
cfg.DATASETS.TEST = ("my_dataset_test",)
cfg.DATASETS.VAL = ("my_dataset_val",)
cfg.DATALOADER.NUM_WORKERS = 2
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-InstanceSegmentation/mask_rcnn_R_50_FPN_3x.yaml") # Let training
cfg.SOLVERIMS_PER_BATCH = 4 # Adjust this parameter (Batch Size)
cfg.SOLVER.BASE_LR = 0.0025 # Adjust this parameter (Learning Rate) initial=0.00025
cfg.SOLVER.MAX_ITER = 500 # 1000 iterations does well for this dataset
cfg.SOLVER.STEPS = [] # do not decay learning rate
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512 # Default is 512
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 2 # We have 2 classes ["Tubes", "Spheres"].

os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = DefaultTrainer(cfg) #Create an instance of of DefaultTrainer with the given configuration
trainer.resume_or_load(resume=False) #Load a pretrained model if available (resume training) or start training from scr
[07/30 14:44:56 d2.engine.defaults]: Model:
GeneralizedRCNN(
    (backbone): FPN(
        (fpn_lateral2): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (fpn_output2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (fpn_lateral3): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
```

```
(fpn_output3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(fpn_lateral4): Conv2d(1024, 256, kernel_size=(1, 1), stride=(1, 1))
(fpn_output4): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(fpn_lateral5): Conv2d(2048, 256, kernel_size=(1, 1), stride=(1, 1))
(fpn_output5): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(top_block): LastLevelMaxPool()
(bottom_up): ResNet(
    (stem): BasicStem(
        (conv1): Conv2d(
            3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False
            (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
        )
    )
    (res2): Sequential(
        (0): BottleneckBlock(
            (shortcut): Conv2d(
                64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
                (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
            )
            (conv1): Conv2d(
                64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
                (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
            )
            (conv2): Conv2d(
                64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
                (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
            )
            (conv3): Conv2d(
                64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False
                (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)
            )
        )
        (1): BottleneckBlock(
            (conv1): Conv2d(
                256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False
                (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
            )
            (conv2): Conv2d(
                64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
                (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)
            )
            (conv3): Conv2d(
```

```
\n        \n        \n        64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False\n        (norm): FrozenBatchNorm2d(num_features=256, eps=1e-05)\n    )\n)\n(2): BottleneckBlock(\n    (conv1): Conv2d(\n        256, 64, kernel_size=(1, 1), stride=(1, 1), bias=False\n        (norm): FrozenBatchNorm2d(num_features=64, eps=1e-05)\n    )\n    (conv2): Conv2d(\n
```

## ▼ Train the model

```
trainer.train()
```

```
[07/30 14:45:05 d2.engine.train_loop]: Starting training from iteration 0\n/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork() was called. os.fork() is incompatible with self.pid = os.fork()\n[07/30 14:45:31 d2.utils.events]: eta: 0:08:50 iter: 19 total_loss: 2.979 loss_cls: 1.049 loss_box_reg: 0.5983\n[07/30 14:45:54 d2.utils.events]: eta: 0:08:28 iter: 39 total_loss: 2.158 loss_cls: 0.6302 loss_box_reg: 0.647\n[07/30 14:46:17 d2.utils.events]: eta: 0:08:08 iter: 59 total_loss: 1.832 loss_cls: 0.5043 loss_box_reg: 0.6613\n[07/30 14:46:40 d2.utils.events]: eta: 0:07:55 iter: 79 total_loss: 1.518 loss_cls: 0.3865 loss_box_reg: 0.6077\n[07/30 14:47:04 d2.utils.events]: eta: 0:07:35 iter: 99 total_loss: 1.3 loss_cls: 0.3201 loss_box_reg: 0.5213\n[07/30 14:47:28 d2.utils.events]: eta: 0:07:14 iter: 119 total_loss: 1.078 loss_cls: 0.2729 loss_box_reg: 0.395\n[07/30 14:47:51 d2.utils.events]: eta: 0:06:50 iter: 139 total_loss: 1.043 loss_cls: 0.275 loss_box_reg: 0.3804\n[07/30 14:48:15 d2.utils.events]: eta: 0:06:28 iter: 159 total_loss: 0.8852 loss_cls: 0.2419 loss_box_reg: 0.30\n[07/30 14:48:38 d2.utils.events]: eta: 0:06:05 iter: 179 total_loss: 0.8003 loss_cls: 0.1957 loss_box_reg: 0.28\n[07/30 14:49:02 d2.utils.events]: eta: 0:05:47 iter: 199 total_loss: 0.8074 loss_cls: 0.1947 loss_box_reg: 0.29\n[07/30 14:49:26 d2.utils.events]: eta: 0:05:24 iter: 219 total_loss: 0.7845 loss_cls: 0.1799 loss_box_reg: 0.28\n[07/30 14:49:50 d2.utils.events]: eta: 0:05:01 iter: 239 total_loss: 0.7438 loss_cls: 0.1856 loss_box_reg: 0.26\n[07/30 14:50:14 d2.utils.events]: eta: 0:04:38 iter: 259 total_loss: 0.7317 loss_cls: 0.1753 loss_box_reg: 0.25\n[07/30 14:50:37 d2.utils.events]: eta: 0:04:15 iter: 279 total_loss: 0.7169 loss_cls: 0.1726 loss_box_reg: 0.25\n[07/30 14:51:02 d2.utils.events]: eta: 0:03:52 iter: 299 total_loss: 0.6586 loss_cls: 0.1443 loss_box_reg: 0.23\n[07/30 14:51:26 d2.utils.events]: eta: 0:03:29 iter: 319 total_loss: 0.7388 loss_cls: 0.1522 loss_box_reg: 0.26\n[07/30 14:51:50 d2.utils.events]: eta: 0:03:06 iter: 339 total_loss: 0.6568 loss_cls: 0.1634 loss_box_reg: 0.23\n[07/30 14:52:13 d2.utils.events]: eta: 0:02:43 iter: 359 total_loss: 0.6692 loss_cls: 0.1609 loss_box_reg: 0.23\n[07/30 14:52:38 d2.utils.events]: eta: 0:02:20 iter: 379 total_loss: 0.6814 loss_cls: 0.1417 loss_box_reg: 0.23\n[07/30 14:53:03 d2.utils.events]: eta: 0:01:56 iter: 399 total_loss: 0.6528 loss_cls: 0.13 loss_box_reg: 0.2398\n[07/30 14:53:29 d2.utils.events]: eta: 0:01:33 iter: 419 total_loss: 0.6574 loss_cls: 0.1397 loss_box_reg: 0.22
```

```
[07/30 14:53:54 d2.utils.events]: eta: 0:01:10 iter: 439 total_loss: 0.646 loss_cls: 0.1369 loss_box_reg: 0.229
[07/30 14:54:17 d2.utils.events]: eta: 0:00:46 iter: 459 total_loss: 0.6063 loss_cls: 0.1235 loss_box_reg: 0.21
[07/30 14:54:42 d2.utils.events]: eta: 0:00:23 iter: 479 total_loss: 0.65 loss_cls: 0.1444 loss_box_reg: 0.2269
[07/30 14:55:13 d2.utils.events]: eta: 0:00:00 iter: 499 total_loss: 0.6354 loss_cls: 0.1275 loss_box_reg: 0.22
[07/30 14:55:13 d2.engine.hooks]: Overall training speed: 498 iterations in 0:09:54 (1.1932 s / it)
[07/30 14:55:13 d2.engine.hooks]: Total training time: 0:10:02 (0:00:08 on hooks)
WARNING [07/30 14:55:13 d2.data.datasets.coco]:
Category ids in annotations are not in [1, #categories]! We'll apply a mapping for you.

[07/30 14:55:13 d2.data.datasets.coco]: Loaded 11 images in COCO format from ./data/annotations/test_annotations.json
[07/30 14:55:13 d2.data.build]: Distribution of instances among all 2 categories:
| category | #instances | category | #instances |
|:-----:|:-----:|:-----:|:-----|
| Rods     | 518       | Spheres   | 0
|          |            |           |
| total    | 518       |           |
[07/30 14:55:13 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used in inference: [ResizeShortestEdge(short_...
[07/30 14:55:13 d2.data.common]: Serializing the dataset using: <class 'detectron2.data.common._TorchSerializedList'...
[07/30 14:55:13 d2.data.common]: Serializing 11 elements to byte tensors and concatenating them all ...
[07/30 14:55:13 d2.data.common]: Serialized dataset takes 0.28 MiB
WARNING [07/30 14:55:13 d2.engine.defaults]: No evaluator found. Use `DefaultTrainer.test(evaluators=)` , or implemen
```

```
# Close the TensorBoard writer
writer.close()
```

## ▼ Save the Configuration

```
import yaml
import datetime

# Create a configuration file with timestamp
timestamp = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")
model_path = "./models/Detectron2_Models/"
yaml_file_name = f"config_{timestamp}.yaml"
config_yaml_path = os.path.join(model_path, yaml_file_name)
with open(config_yaml_path, 'w') as file:
    ...yaml dump here...

```
yaml
```


```

```
yaml.dump(cfg, title)

model_weight_filename = f'model_weights_{timestamp}.pth'
model_weight_path = os.path.join(model_path, model_weight_filename)
torch.save(trainer.model.state_dict(), model_weight_path)
```

## ✓ Evaluation of trained model

```
def load_model(config_file_path, model_weights_path, score_thresh=0.6):
    """
    Load the Detectron2 model with the given configuration and weights.

    Args:
        config_file_path (str): Path to the configuration file.
        model_weights_path (str): Path to the model weights file.
        score_thresh (float): Score threshold for the model. Default is 0.6.

    Returns:
        DefaultPredictor: The predictor with the loaded model.
    """

    cfg = get_cfg()
    cfg.merge_from_file(config_file_path)
    cfg.MODEL.WEIGHTS = model_weights_path
    cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = score_thresh
    predictor = DefaultPredictor(cfg)

    return predictor
```

Uncomment the below code to load a previously trained model.

```
#config_file_path = "./models/Detectron2_Models/config.yaml"
#model_weights_path = "./models/Detectron2_Models/model_final.pth"
#load_model(config_file_path, model_weights_path, score_thresh=0.6)
```

## ▼ TensorBoard

```
%load_ext tensorboard  
%tensorboard --logdir ./output/logs #./models/Detectron2_Models
```

The tensorboard extension is already loaded. To reload it, use:  
  %reload\_ext tensorboard

TensorBoard

INACTIVE

**No dashboards are active for the current data set.**

Probable causes:

- You haven't written any data to your event files.
- TensorBoard can't find your event files.

If you're new to using TensorBoard, and want to find out how to add data and set up your event files, check out the [README](#) and perhaps the [TensorBoard tutorial](#).

If you think TensorBoard is configured properly, please see [the section of the README devoted to missing data problems](#) and consider filing an issue on GitHub.

*Last reload: Jul 30, 2024, 11:14:11 AM*

*Log directory: ./output/logs*

```
# For evaluation
from detectron2.evaluation import COCOEvaluator, inference_on_dataset
from detectron2.data import build_detection_test_loader

# Evaluate the model
cfg.MODEL.WEIGHTS = os.path.join(cfg.OUTPUT_DIR, "model_final.pth")
cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.5
predictor = DefaultPredictor(cfg)

evaluator = COCOEvaluator("my_dataset_val", output_dir=".output")
val_loader = build_detection_test_loader(cfg, "my_dataset_val")
print(inference_on_dataset(predictor.model, val_loader, evaluator))

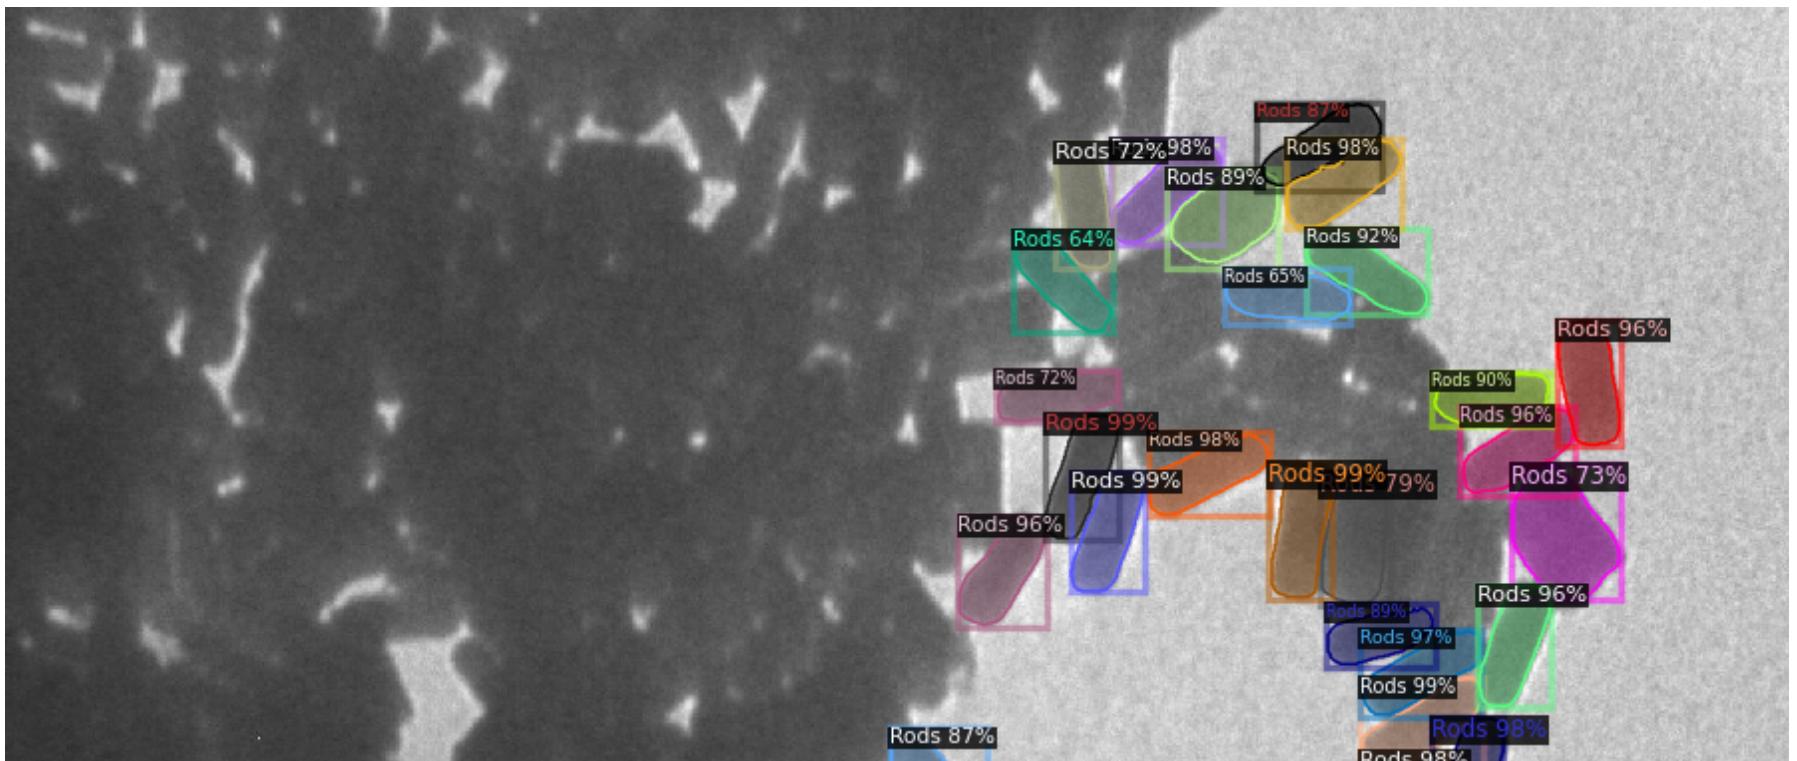
[07/30 04:46:04 d2.checkpoint.detection_checkpoint]: [DetectionCheckpointer] Loading from ./models/Detectron2_Models
[07/30 04:46:04 d2.evaluation.coco_evaluation]: Fast COCO eval is not built. Falling back to official COCO eval.
WARNING [07/30 04:46:04 d2.data.datasets.coco]:
Category ids in annotations are not in [1, #categories]! We'll apply a mapping for you.

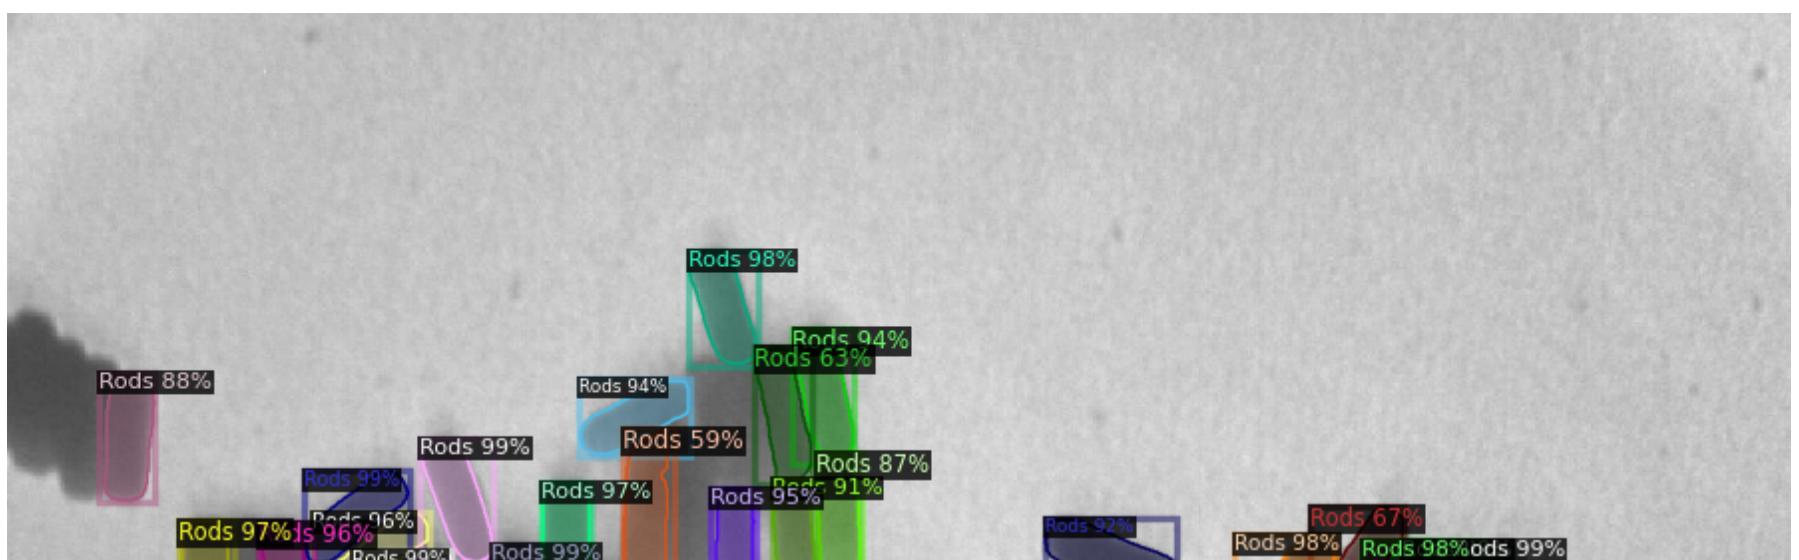
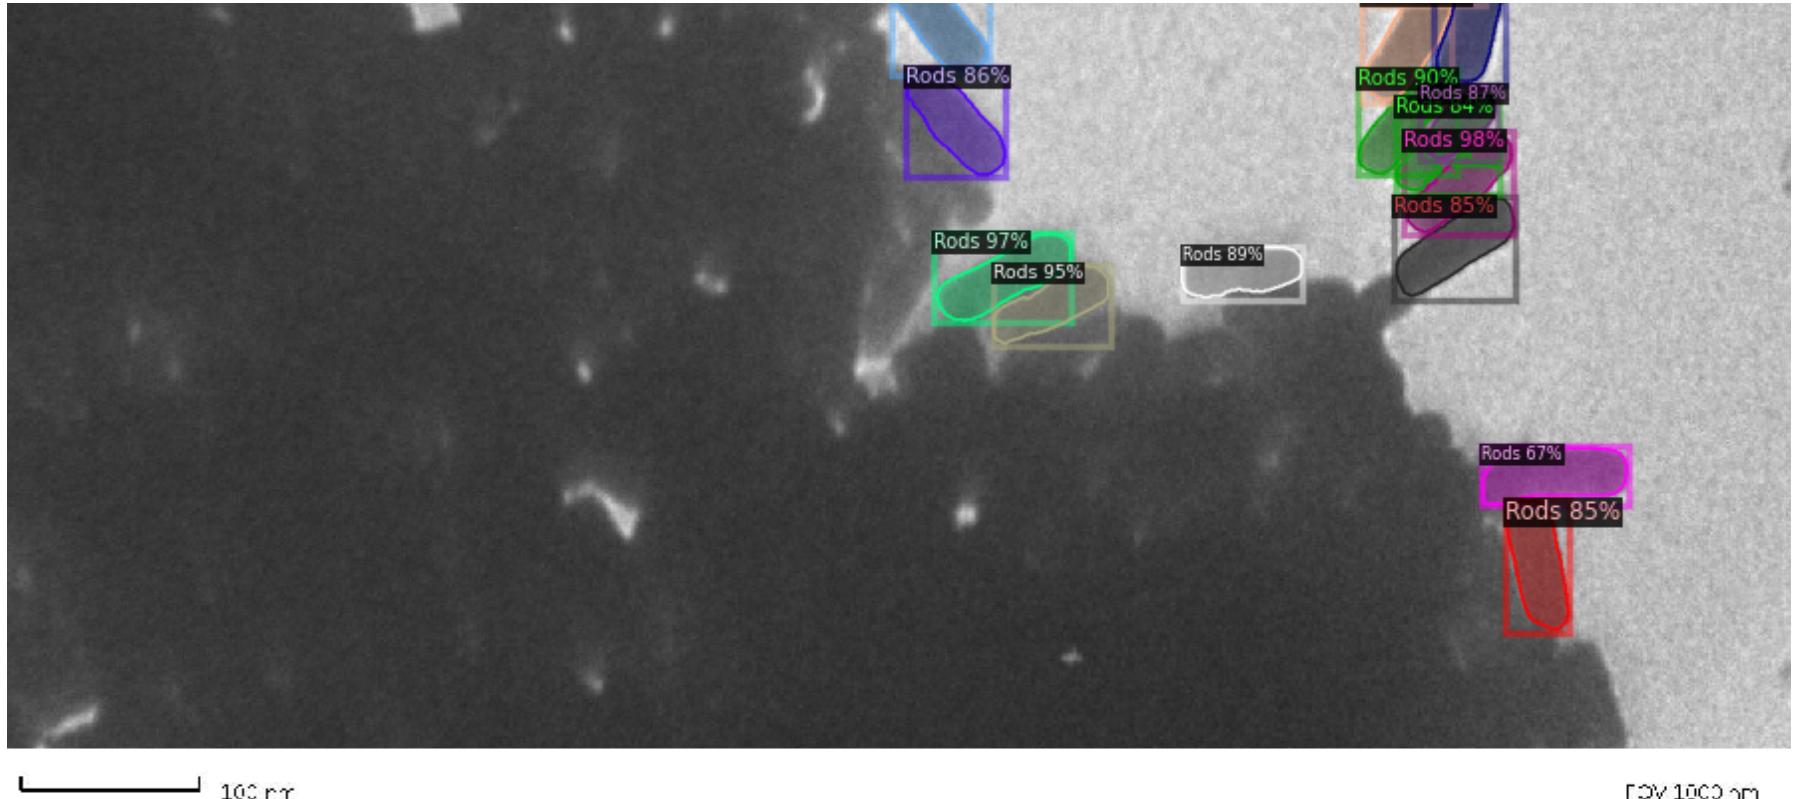
[07/30 04:46:04 d2.data.datasets.coco]: Loaded 21 images in COCO format from ./data/annotations/val_annotations.json
[07/30 04:46:04 d2.data.dataset_mapper]: [DatasetMapper] Augmentations used in inference: [ResizeShortestEdge(short_
[07/30 04:46:04 d2.data.common]: Serializing the dataset using: <class 'detectron2.data.common._TorchSerializedList'>
[07/30 04:46:04 d2.data.common]: Serializing 21 elements to byte tensors and concatenating them all ...
[07/30 04:46:04 d2.data.common]: Serialized dataset takes 0.91 MiB
[07/30 04:46:04 d2.evaluation.evaluator]: Start inference on 21 batches
/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork() was called. os.fork() is incompatibl
```

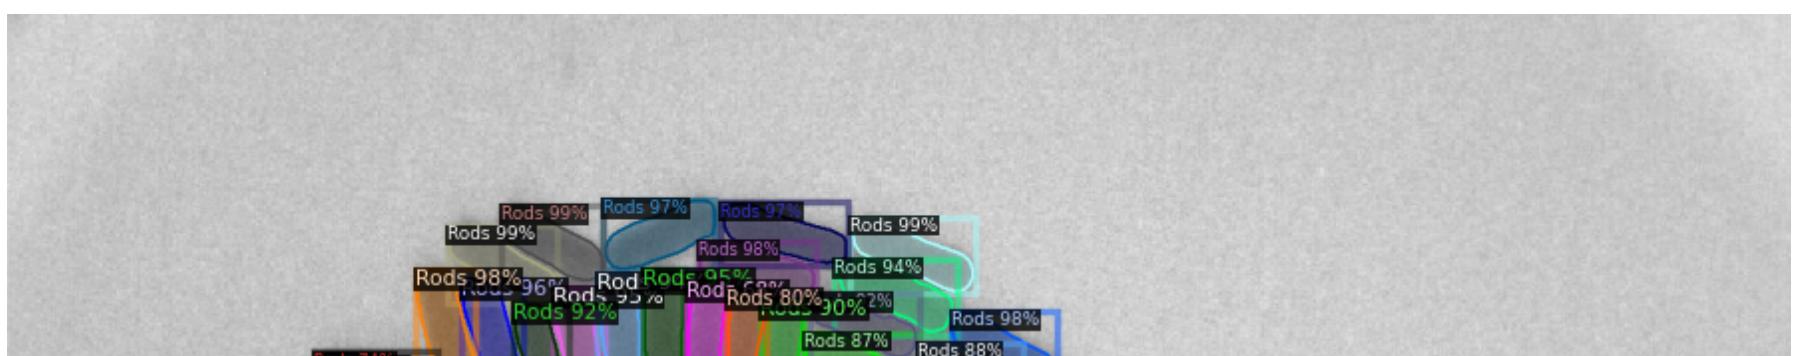
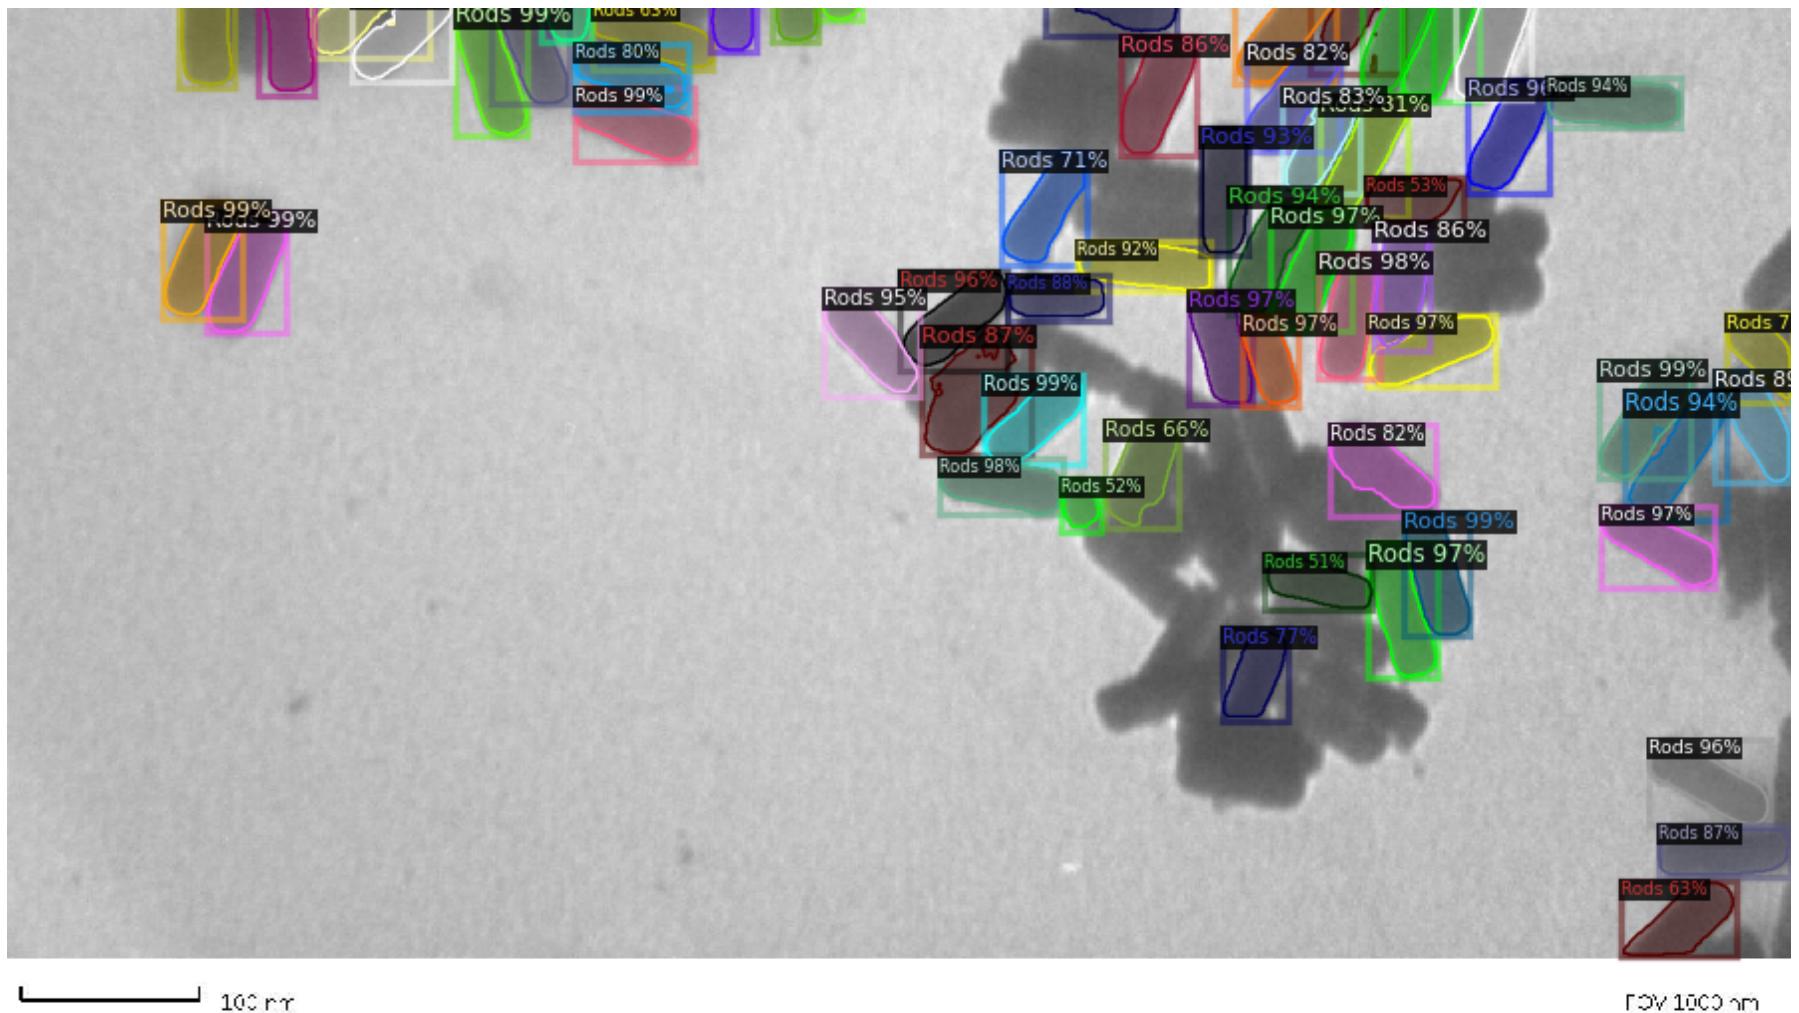
```
self.pid = os.fork()
[07/30 04:46:09 d2.evaluation.evaluator]: Inference done 11/21. Dataloading: 0.0017 s/iter. Inference: 0.1326 s/iter
[07/30 04:46:15 d2.evaluation.evaluator]: Inference done 19/21. Dataloading: 0.0042 s/iter. Inference: 0.1714 s/iter
[07/30 04:46:16 d2.evaluation.evaluator]: Total inference time: 0:00:08.844365 (0.552773 s / iter per device, on 1 d
[07/30 04:46:16 d2.evaluation.evaluator]: Total inference pure compute time: 0:00:02 (0.168142 s / iter per device,
[07/30 04:46:16 d2.evaluation.coco_evaluation]: Preparing results for COCO format ...
[07/30 04:46:16 d2.evaluation.coco_evaluation]: Saving results to ./output/coco_instances_results.json
[07/30 04:46:16 d2.evaluation.coco_evaluation]: Evaluating predictions with official COCO API...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=2.63s).
Accumulating evaluation results...
DONE (t=0.02s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.562
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.716
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.694
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.483
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.567
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.010
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.092
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.607
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.507
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.624
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = -1.000
[07/30 04:46:18 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP      | AP50    | AP75    | APs     | APm     | AP1     |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 56.171 | 71.577 | 69.361 | 48.273 | 56.687 | nan     |
[07/30 04:46:18 d2.evaluation.coco_evaluation]: Some metrics cannot be computed and is shown as NaN.
[07/30 04:46:18 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP      | category | AP      |
|:-----:|:-----:|:-----:|:-----:|
| Rods    | 49.264 | Spheres  | 63.079 |
Loading and preparing results...
DONE (t=0.01s)
creating index...
index created!
```

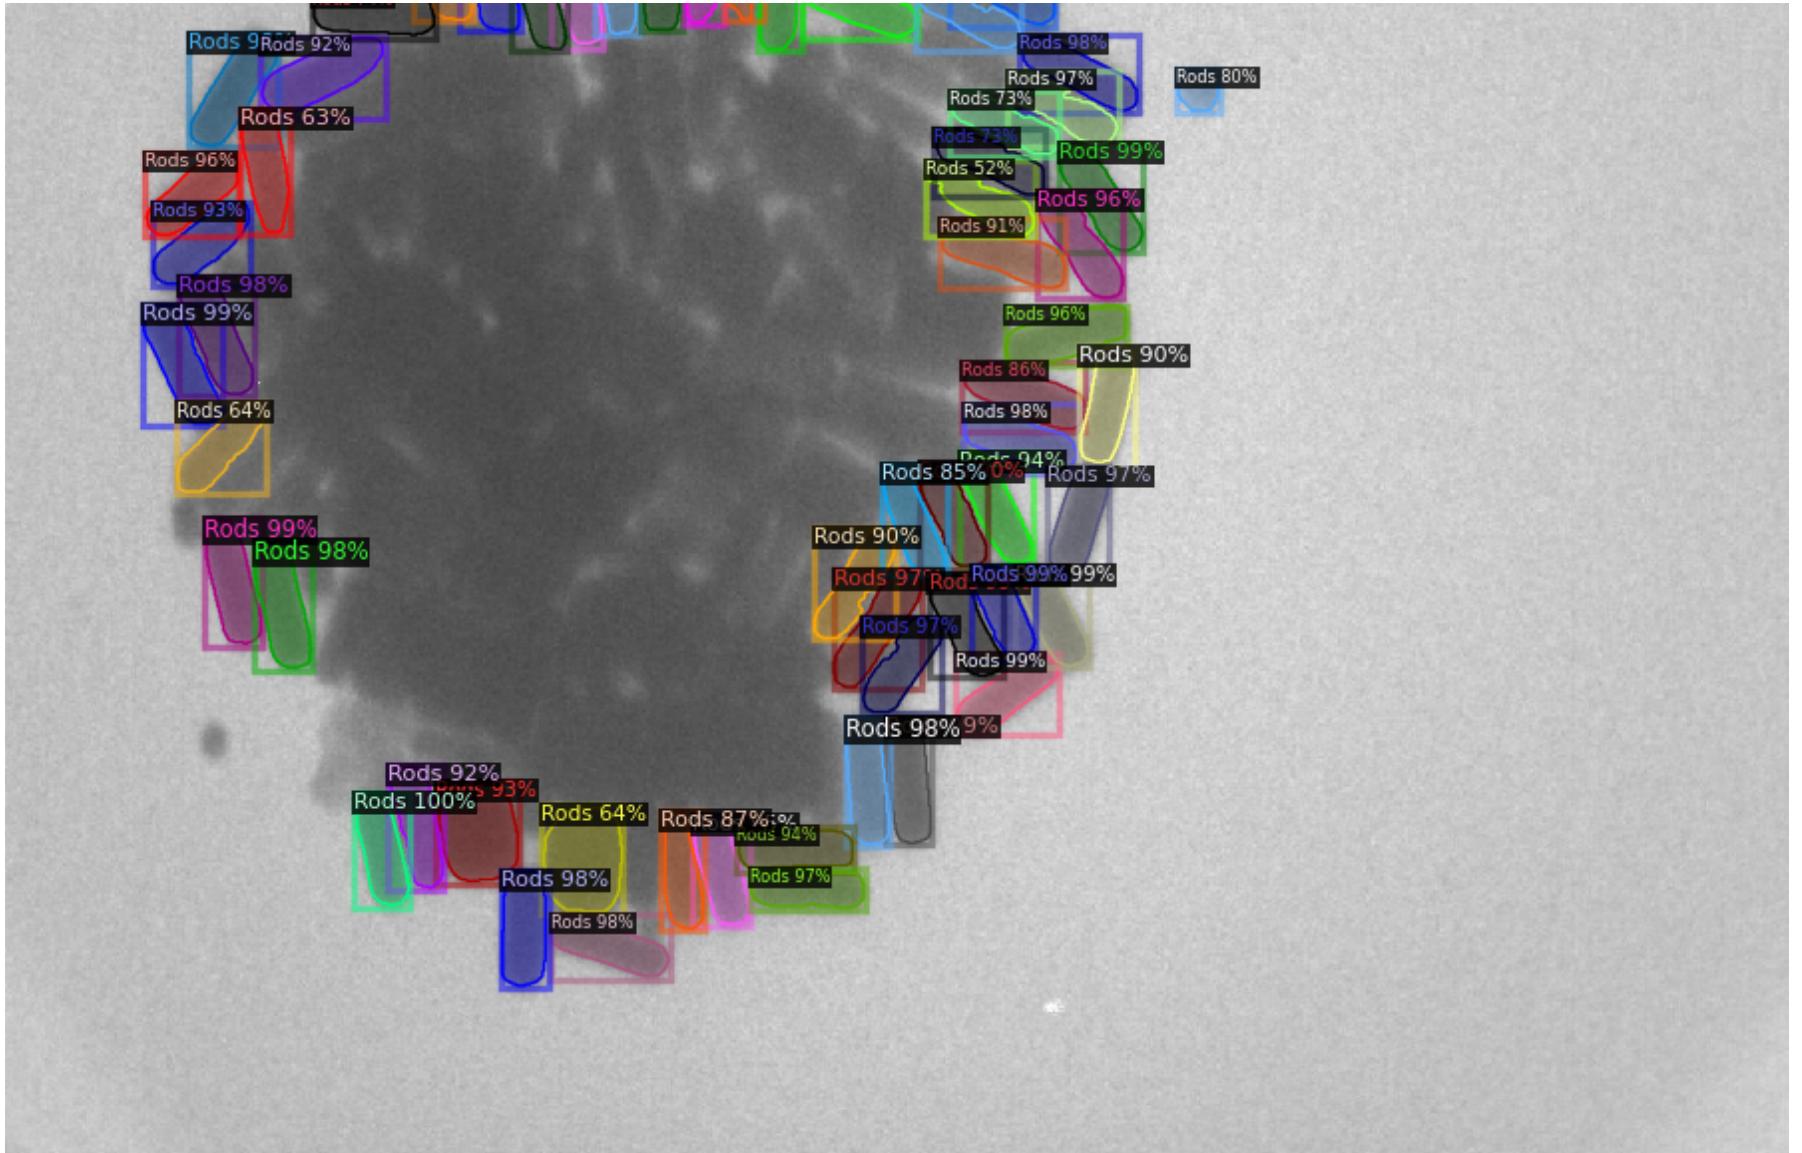
```
index created:  
Running per image evaluation...  
Evaluate annotation type *segm*  
DONE (t=2.62s).  
Accumulating evaluation results...
```

```
for d in random.sample(val_dataset_dicts, 4):      #select number of images for display  
    im = cv2.imread(d["file_name"])  
    outputs = predictor(im)  
    v = Visualizer(im[:, :, ::-1],  
                   metadata=val_metadata,  
                   scale=0.7,  
                   instance_mode=ColorMode.IMAGE_BW    # remove the colors of unsegmented pixels. This option is only ava  
)  
    out = v.draw_instance_predictions(outputs["instances"].to("cpu"))  
    cv2.imshow(out.get_image()[:, :, ::-1])  
    output_filepath = os.path.join("./output/predictions/images/", os.path.basename(d["file_name"]))  
    cv2.imwrite(output_filepath, out.get_image()[:, :, ::-1])
```



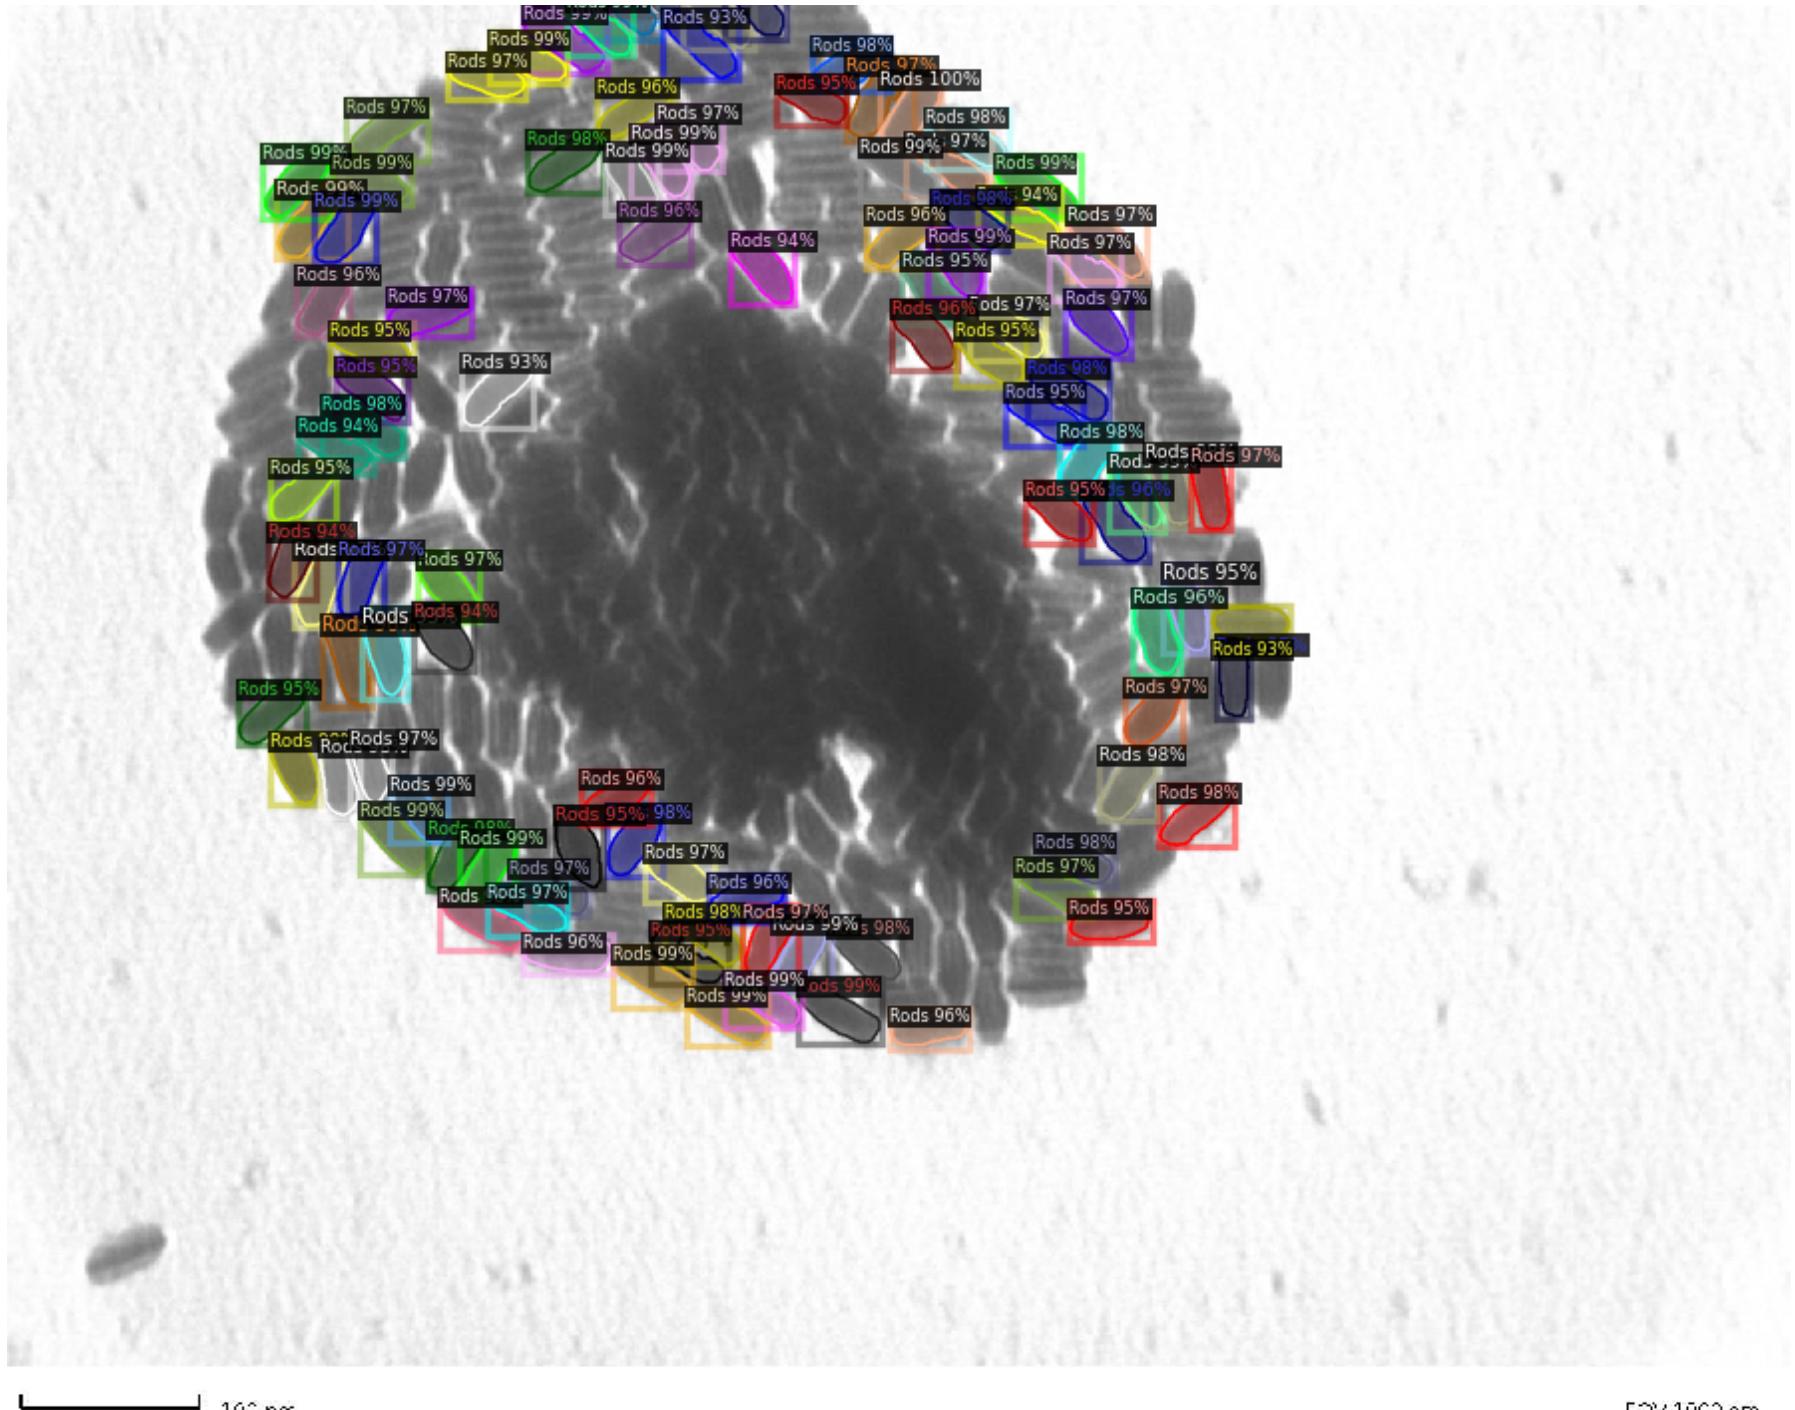






100 μm

COV 1000 nm



Double-click (or enter) to edit

```
import pandas as pd

def run_inference_and_save_results(cfg, image_dir, output_csv):
    predictor = DefaultPredictor(cfg)
    results = []

    for filename in os.listdir(image_dir):
        if filename.endswith(".jpg") or filename.endswith(".png"):
            file_path = os.path.join(image_dir, filename)
            image = cv2.imread(file_path)
            outputs = predictor(image)

            pred_classes = outputs["instances"].pred_classes.cpu().numpy()
            pred_masks = outputs["instances"].pred_masks.cpu().numpy()
            unique, counts = np.unique(pred_classes, return_counts=True)
            class_counts = dict(zip(unique, counts))

            num_rods = class_counts.get(0, 0) # Assuming 'Rod' is class 0
            num_spheres = class_counts.get(1, 0) # Assuming 'Sphere' is class 1

            rod_area = []
            sphere_area = []

            for i, cls in enumerate(pred_classes):
                area = np.sum(pred_masks[i])
                if cls == 0: # Rods
                    rod_area.append(area)
                elif cls == 1: # Spheres
                    sphere_area.append(area)

            results.append({
```

```
        "filename": filename,
        "number_of_rod": num_rod,
        "number_of_spheres": num_spheres,
        "rod_area": rod_area,
        "sphere_area": sphere_area
    })

df = pd.DataFrame(results)
df.to_csv(output_csv, index=False)
print(f"Results saved to {output_csv}")
```

## ▼ Save Data to csv file

```
# Run inference and save results
image_dir = "./data/images/evaluate/"
output_csv = "./output/results/output.csv"
run_inference_and_save_results(cfg, image_dir, output_csv)

[07/30 04:46:29 d2.checkpoint.detection_checkpoint]: [DetectionCheckpointer] Loading from ./models/Detectron2_Models
Results saved to ./output/results/output.csv
```

